


# Database Technology

## Exercise 6: Review



## Q1: I/O costs of each access path

Consider the following relation  $R(x,y,z)$ . The table  $R$  has 10000 tuples which are stored on 600 blocks. There are individual index available on each attribute  $x$ ,  $y$ , and  $z$  but only the index on  $y$  is clustered.

The attribute domain cardinalities are as follows:

- $V(R, x) = 100$
- $V(R, y) = 20$
- $V(R, z) = 200$

You may assume that the cost to access index blocks is negligible (i.e., it should not contribute to your answer).

Full table scan with filter on all attributes	numOfBlocks = 600
$\sigma_{x=1}(R)$ using index scan on $x$	$\text{Ceil}(\text{numOfTuples}/\text{domain card.}) = \text{Ceil}(10000/100) = 100$
$\sigma_{y=2}(R)$ using index scan on $y$	$\text{Ceil}(\text{numOfBlocks}/\text{domain card.}) = \text{Ceil}(600/20) = 30$
$\sigma_{z=3}(R)$ using index scan on $z$	$\text{Ceil}(\text{numOfTuples}/\text{domain card.}) = \text{Ceil}(10000/200) = 50$

## Q2: Blockwise nested loop join is a pipeline breaker?

True

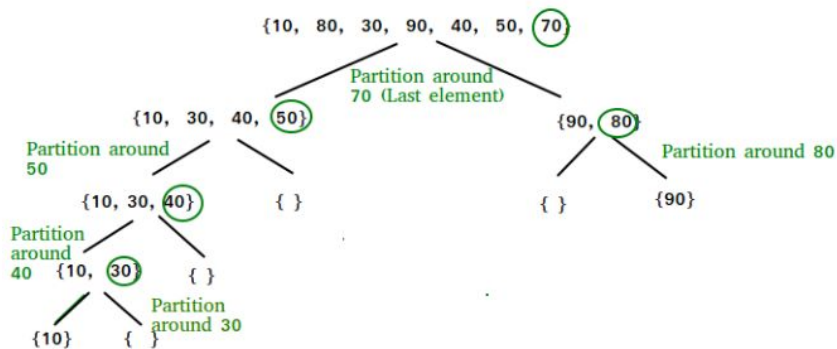
```
FOR EACH chunk of M-1 blocks of S DO BEGIN
  read blocks into main memory;
  organize records into efficient data structure;
  FOR EACH block b of R DO BEGIN
    read b into main memory;
    FOR EACH tuple t of b DO BEGIN
      find records from S in main memory that join;
      output join results;
    END;
  END;
END;
```



Pipeline Breaker  
because S is to  
be materialized  
in memory

## Q2: Quick sort is not a pipeline breaker?

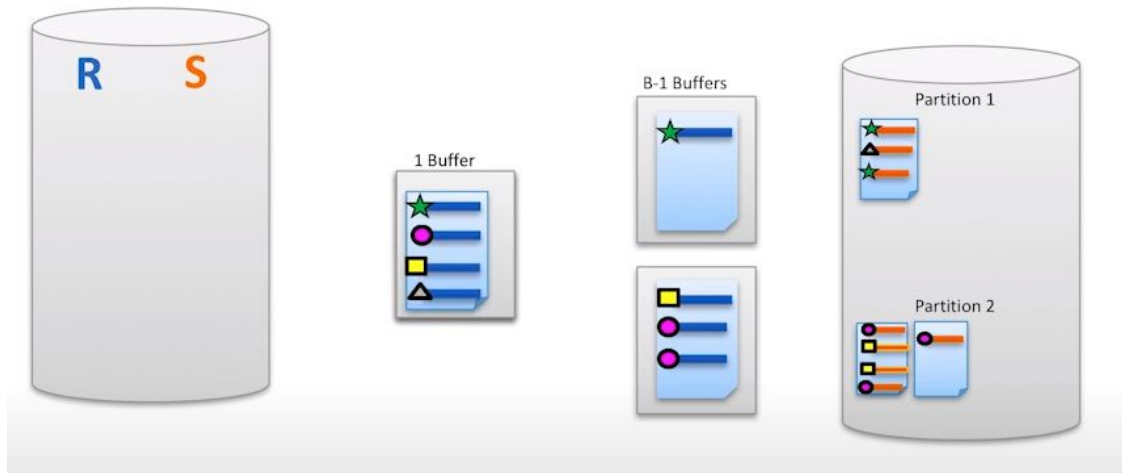
False



Need whole data before performing sorting. (any sorting algorithm in-general will be a pipeline breaker)

## Q2: Grace hash join is not a pipeline breaker?


False



1. Partition
2. Build Hash Table

### Q3: Identify the pipeline breaker for both tables.

Given are two tables  $R(t_1, s_1)$  other  $S(s_1, s_2)$ . The table  $R$  is clustered on  $R.s_1$ . The index on  $S.s_1$  is a B+ tree. The tables are joined on the attributes  $s_1$  using the physical plan below.



<b>Table R</b>	No Pipeline Breaker
<b>Table S</b>	No Pipeline Breaker

Q4: Selection and projection do not have a constant memory requirement.

False, Memory Cost = 1



Q5: The little relation should be in the outside loop for better performance of block based NLJ?

True

```
FOR EACH chunk of M-1 blocks of S DO BEGIN
  read blocks into main memory;
  organize records into efficient data structure;
  FOR EACH block b of R DO BEGIN
    read b into main memory;
    FOR EACH tuple t of b DO BEGIN
      find records from S in main memory that join;
      output join results;
    END;
  END;
END;
```



Outer loop will be smaller and we can store everything in memory and thus it will have less pipeline breakers

Q6: For a block-based nested loop join (using the **approximation** formula), what is the expected Disk IO cost?

Given are two relations  $R$  and  $S$ , with block size of  $B(R) = 200$  and  $B(S) = 300$ , tuples sizes of  $T(R) = 2000$  and  $T(S) = 2400$ , and a memory size of  $M = 100$ .

Assume that the relations are unclustered.

$$T(R) * T(S) / M = 2000 * 2400 / 100 = 48000$$

### Operations on Nonclustered Data

All our calculations regarding the number of disk I/O's required for an operation are predicated on the assumption that the operand relations are clustered. In the (typically rare) event that an operand  $R$  is not clustered, then it may take us  $T(R)$  disk I/O's, rather than  $B(R)$  disk I/O's to read all the tuples of  $R$ . Note, however, that any relation that is the result of an operator may always be assumed clustered, since we have no reason to store a temporary relation in a nonclustered fashion.

Q7: The IO cost for sort based two pass duplication elimination is  $3B(R)$ .

True

1.  $B(R)$  for **reading** in Phase 1
  2.  $B(R)$  for **writing** list partitions
  3.  $B(R)$  for **re-reading** list partitions
- Together:  $3 \cdot B(R)$

Q8: Duplicate elimination cannot be done using a two pass hash based algorithm because the memory requirement becomes unsustainable in the second pass.

False

Q9: For a simple ***sort-based join***, what is the Disk IO Cost and Main Memory requirement?

Assume that the relations are clustered. Given are two relations R and S, with block size of  $B(R) = 400$  and  $B(S) = 225$ , tuples sizes of  $T(R) = 4400$  and  $T(S) = 4500$ , and a memory size of  $M = 100$

Disk IO Cost	$5(B(R)+B(S))=3125$
Main Memory Requirement?	$Sqr(\max(B(R),B(S))) = Sqr(400) = 20$

Q10: For a simple ***hash-based join***, what is the Disk IO Cost and Main Memory requirement?

Assume that the relations are clustered. Given are two relations R and S, with block size of  $B(R) = 400$  and  $B(S) = 400$ , tuples sizes of  $T(R) = 8000$  and  $T(S) = 8000$ , and a memory size of  $M = 175$

Disk IO Cost	$3(B(R)+B(S))=2400$
Main Memory Requirement?	$\text{Sqr}(B(S)) = \text{Sqr}(400) = 20$